

Tal Ben-Nun

Parallel and Distributed Deep Learning



Where is Deep Learning used?

Digit Recognition



Object Classification
Segmentation

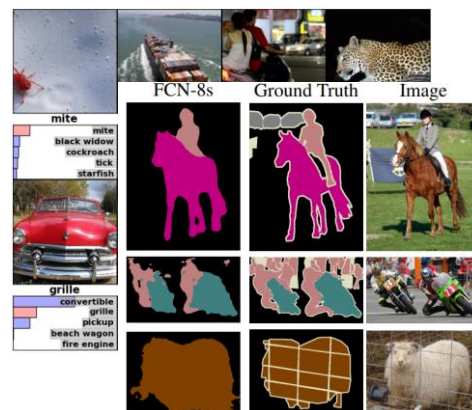




Image Captioning





Example output of the model

Gameplay AI
Translation



Neural Computers
Routing



Why Scale Up?

- Enormous amounts of data
 - MSCOCO: 19 GB
 - ImageNet (1k): 180 GB
 - ImageNet (22k): A few TB
 - Industry: Much larger

- Large neural network architectures
 - 100-200 layers deep today, ~100M-2B parameters


- Faster prototyping
 - Training time: 10s of hours to days (and weeks)



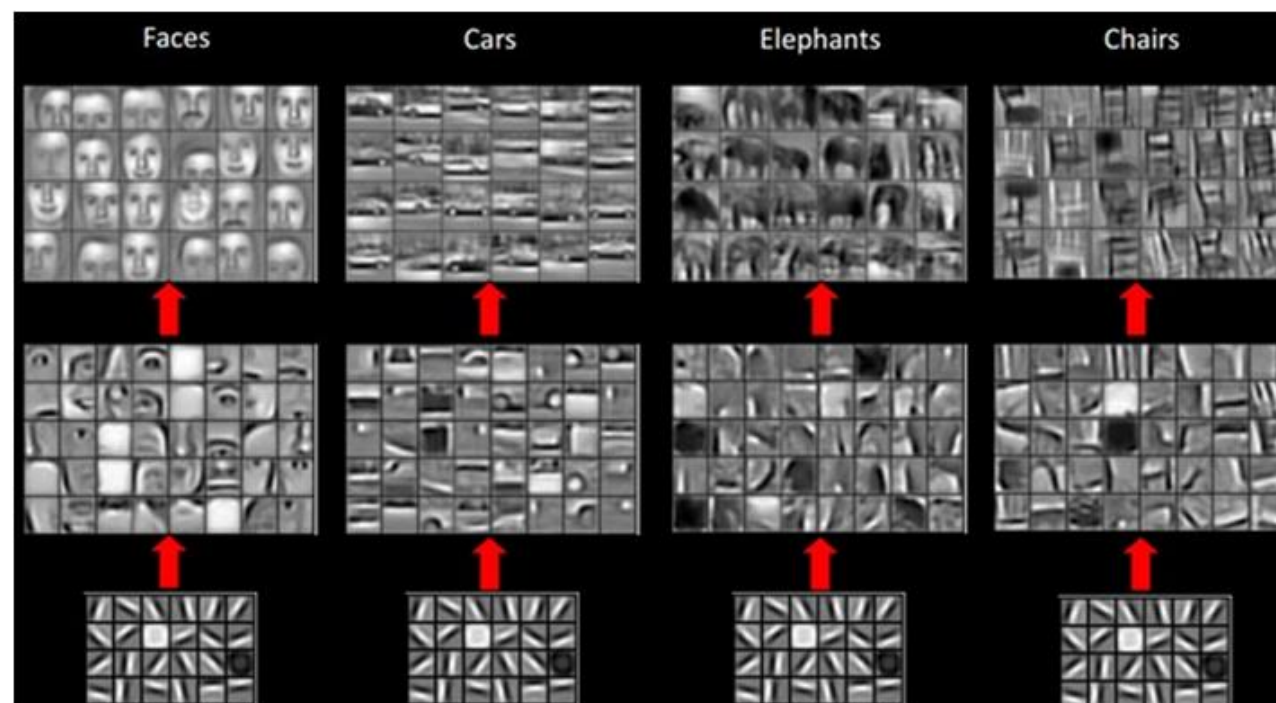
Neural Networks

Neural Networks

- Modeled after the human brain
- CNNs repeatedly perform convolutions and nonlinearity operations




Source: [CS231n](#)



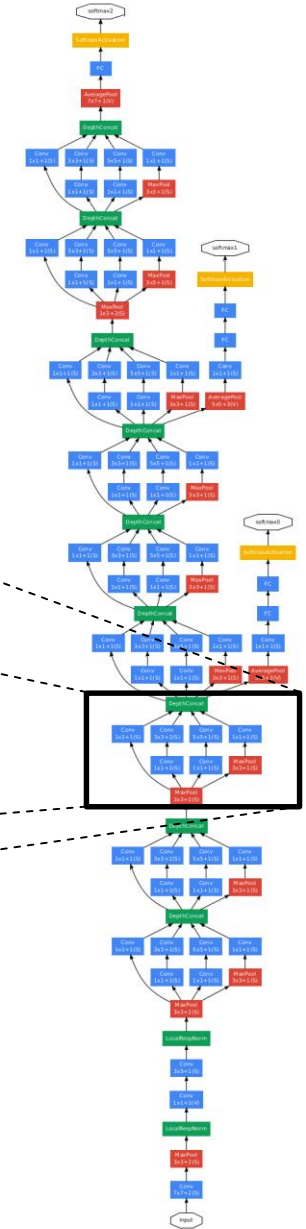
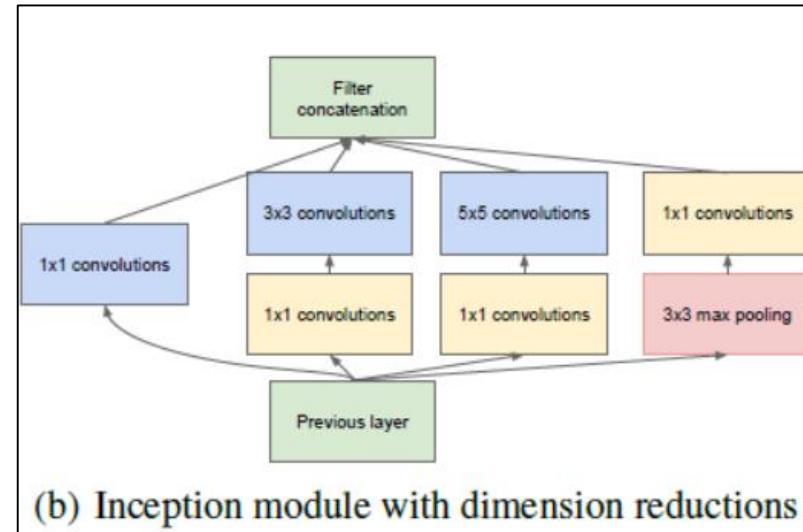
Source: Lee et al. "[Unsupervised Learning of Hierarchical Representations with Convolutional Deep Belief Networks](#)" (CACM 2011)

Simple CNN Architecture





GoogLeNet [Szegedy et al., 2014]

- ~6.8M parameters
- 22 layers deep



- ~2.35M parameters
- 152 layers deep




11

Stochastic Gradient Descent

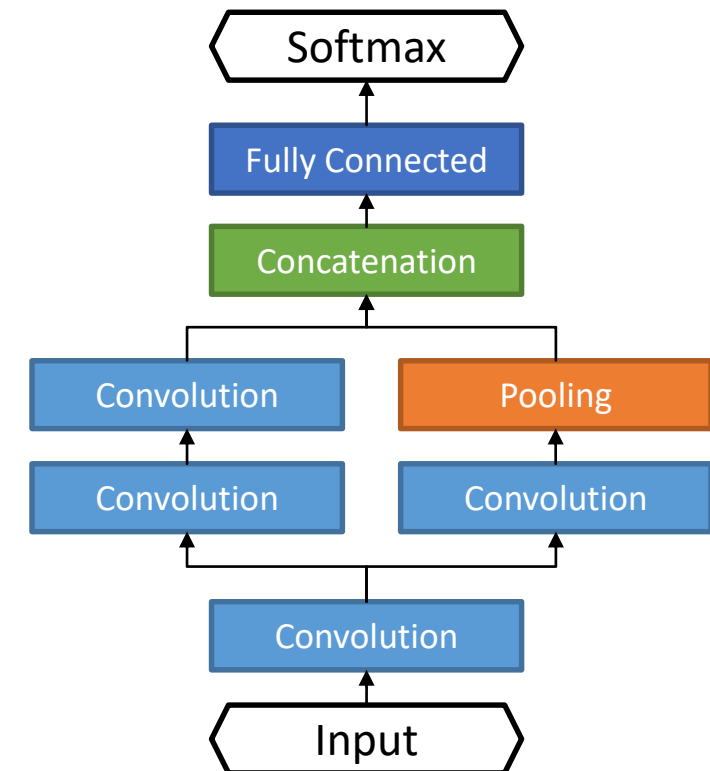
- Gist: Improve network weights using samples from a labeled dataset
- Algorithm:
 - Initialize neural network weights (W_0)
 - For t in iterations:
 - **Sample** b images from dataset (B)
 - **Compute** loss $L_t(W_{t-1}, B)$
 - **Update** weights using gradients and update rule g :
 $W_t = g(W_{t-1}, \nabla L_t(W_{t-1}, B), [\text{hyperparameters} \dots])$
- $\nabla L_t(W, B)$ is an average direction of the gradient over a mini-batch of size b :

$$\nabla L_t(W, B) = \frac{1}{b} \sum_{i=1}^b \nabla \ell(W; (x_i, y_i))$$



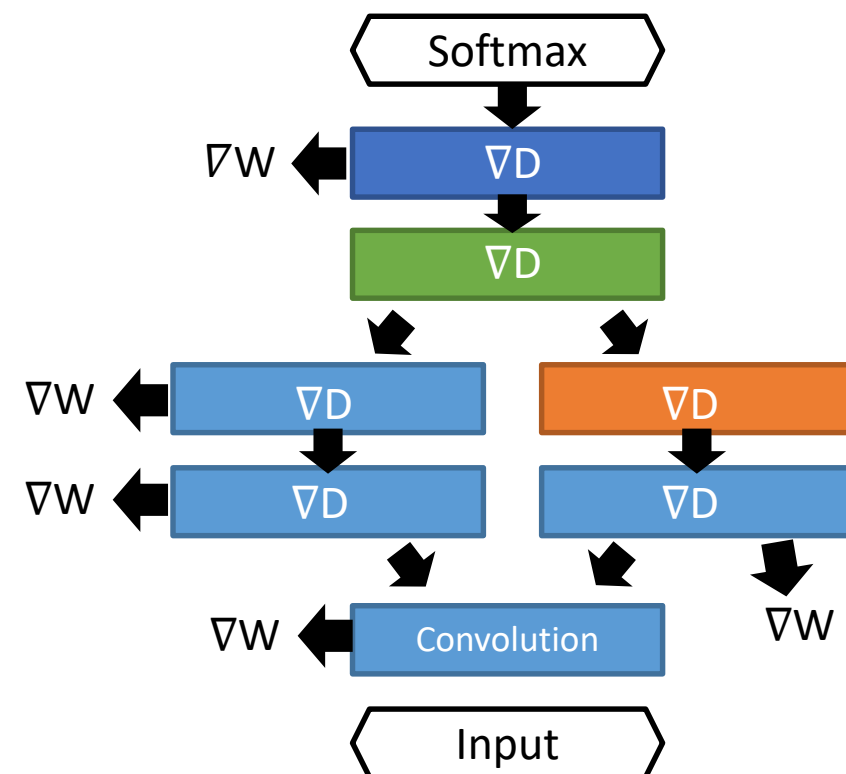
Backpropagation Algorithm

- A CNN is a Directed Acyclic Graph (DAG)
- At each layer in backpropagation, derivatives are estimated w.r.t.:
 - Layer parameters (if necessary)
 - Data (chain rule)



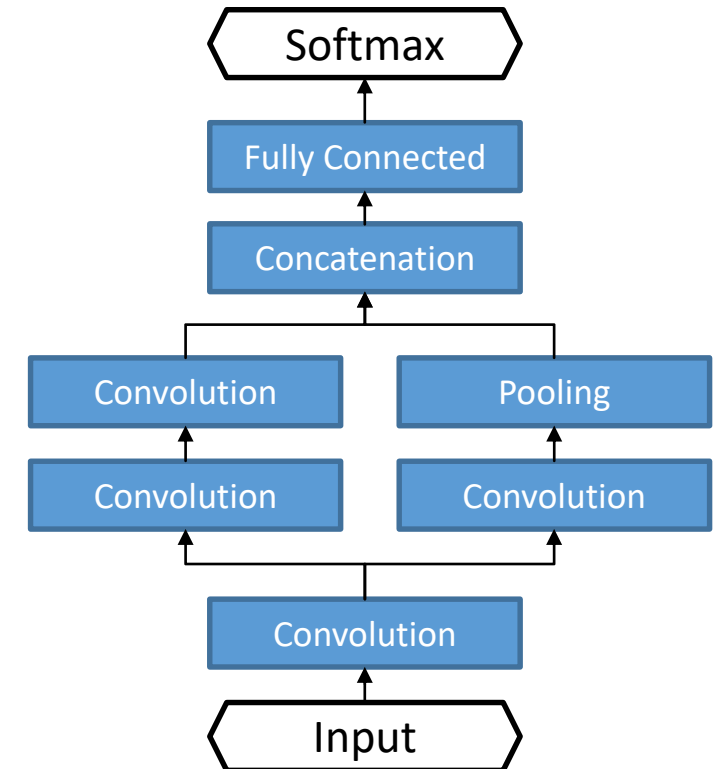
Backpropagation Algorithm

- A CNN is a Directed Acyclic Graph (DAG)
- At each layer in backpropagation, derivatives are estimated w.r.t.:
 - Layer parameters (if necessary)
 - Data (chain rule)
- Additional memory storage required for training:
 - $D+W+\nabla D+\nabla W$



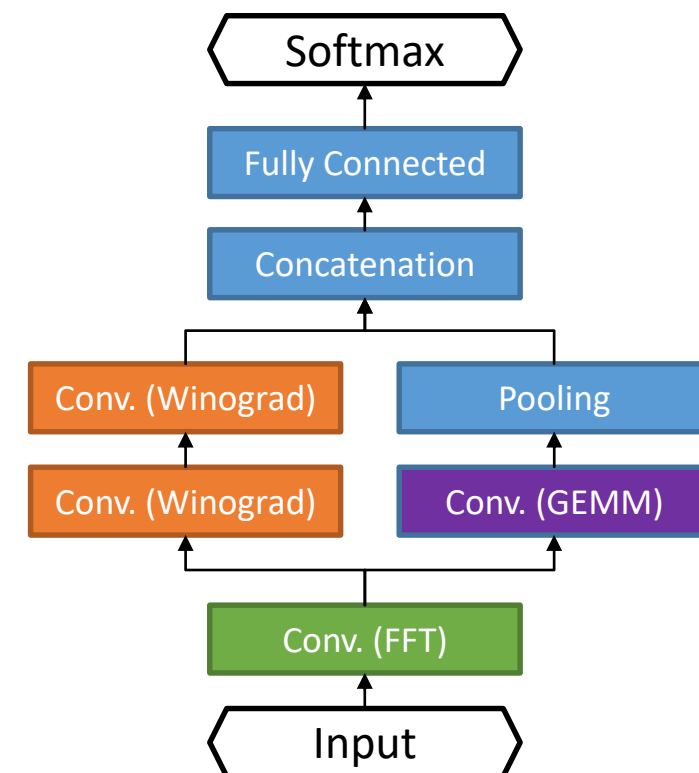
The World of Neural Network Acceleration

- Choice of Algorithm
- Parallelism
- Distributed Computing
- Hardware Architectures



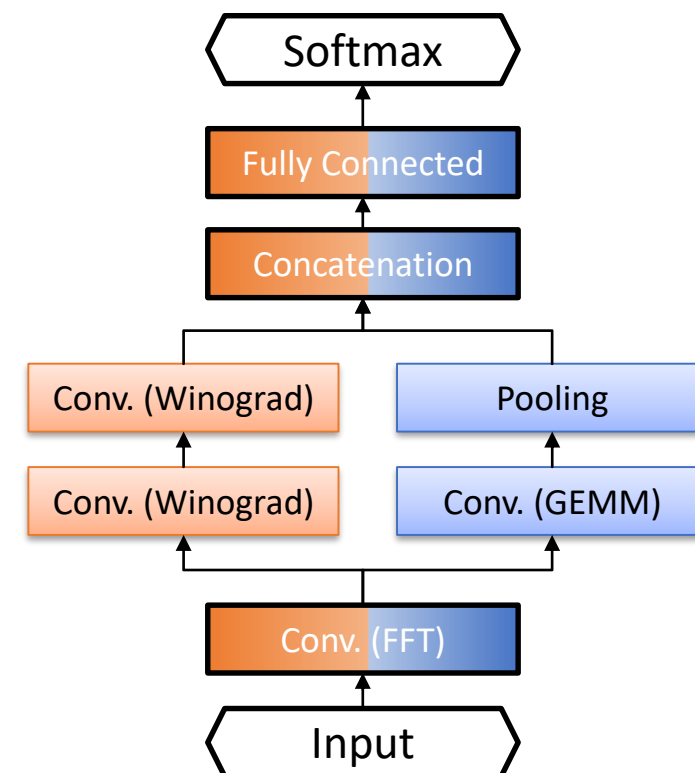
The World of Neural Network Acceleration

- Choice of Algorithm
- Parallelism
- Distributed Computing
- Hardware Architectures




The World of Neural Network Acceleration

- Choice of Algorithm
- **Parallelism**
- Distributed Computing
- Hardware Architectures




The World of Neural Network Acceleration

- Choice of Algorithm
- Parallelism
- **Distributed Computing**
- Hardware Architectures




The World of Neural Network Acceleration

- Choice of Algorithm
- Parallelism
- Distributed Computing
- **Hardware Architectures**



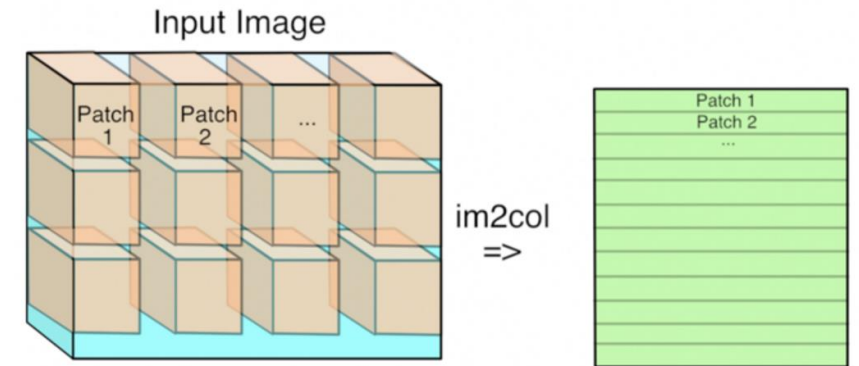
Algorithms




Convolution Algorithms

- Most computationally-intensive layer

$$out(x, y)^{f_o} = \sum_{f_i=0}^{N_{if}} \sum_{k_x=0}^{K_x} \sum_{k_y=0}^{K_y} w_{f_i, f_o}(k_x, k_y) * in(x + k_x, y + k_y)^{f_i}$$



- Can be performed directly, or:
 - Via matrix multiplication (im2col) [Chellapilla et al., 2006]
 - Via Winograd convolution [Lavin and Gray, 2016]
 - In Fourier domain



Source: [Pete Warden's Blog](#)

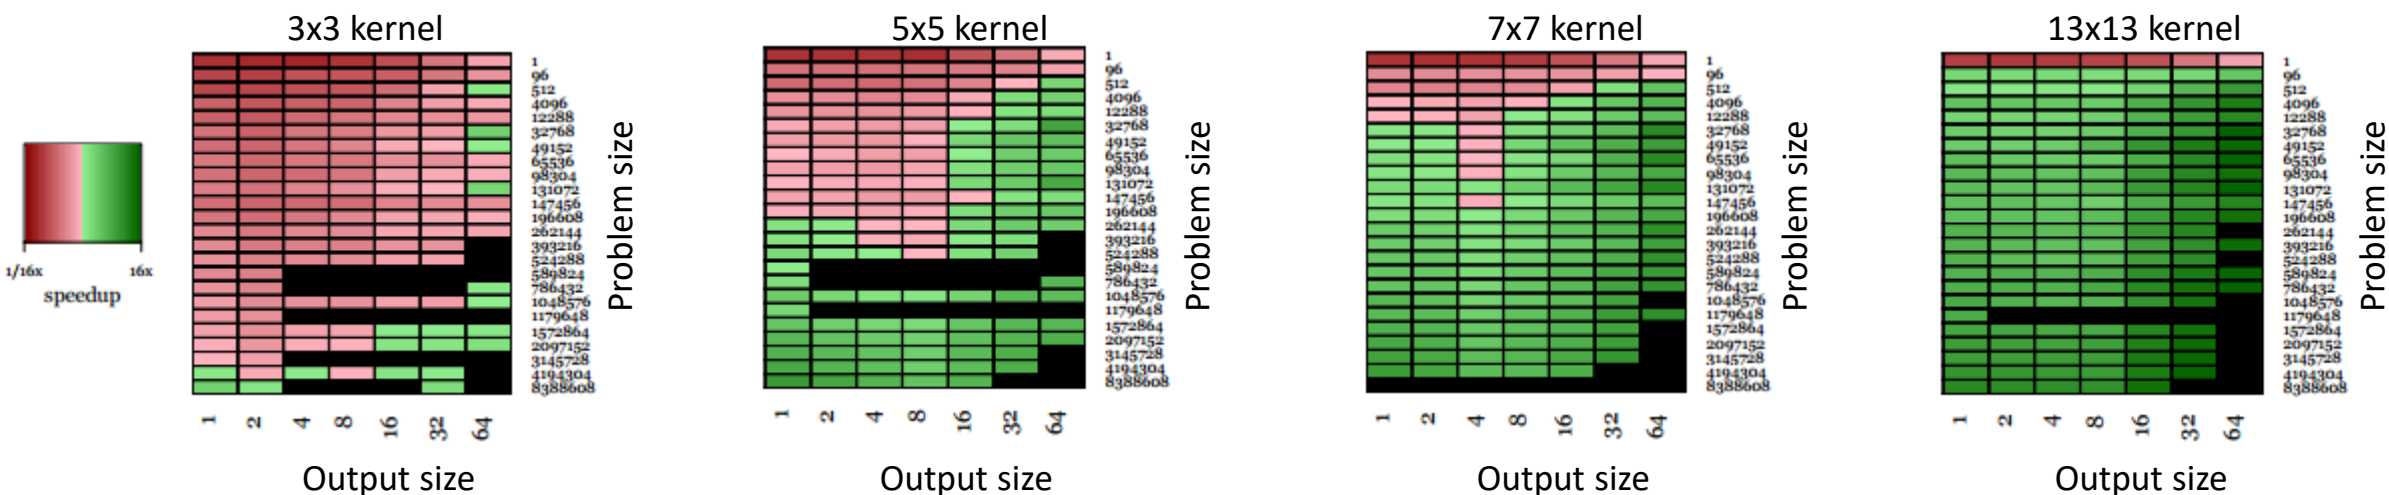
im2col convolution

Convolution in Fourier Domain

- Convolution can be computed using FFT [Mathieu et al., 2014]:

$$y_{(s,j)} = \sum_{i \in f} x_{(s,i)} \star w_{(j,i)} = \sum_{i \in f} \mathcal{F}^{-1} (\mathcal{F}(x_{(s,i)}) \circ \mathcal{F}(w_{(j,i)})^*)$$

- The larger the convolution kernel, the better the performance [Vasilache et al., 2015]:



[Mathieu et al. "Fast Training of Convolutional Networks through FFTs", ICLR 2014](#)

[Vasilache et al. "Fast Convolutional Nets With fbfft: A GPU Performance Evaluation", ICLR 2015](#)

Sacrificing Accuracy for Performance



- Half-precision (16-bit floating point) [Gupta et al., 2015]
 - Memory is stored in 16-bit format
 - Computations are performed in 32-bits
 - Uses Stochastic Rounding:

$$\text{Round}(x, \langle \text{IL}, \text{FL} \rangle) = \begin{cases} \lfloor x \rfloor & \text{w.p. } 1 - \frac{x - \lfloor x \rfloor}{\epsilon} \\ \lfloor x \rfloor + \epsilon & \text{w.p. } \frac{x - \lfloor x \rfloor}{\epsilon} \end{cases}$$

Goal: Preserve $\mathbb{E}(\text{Round}(x, \langle \text{IL}, \text{FL} \rangle)) = x$

Sacrificing Accuracy for Performance


- Results on MNIST with LeNet:




WL=Word Length (bits)

FL=Fractional Length (bits)

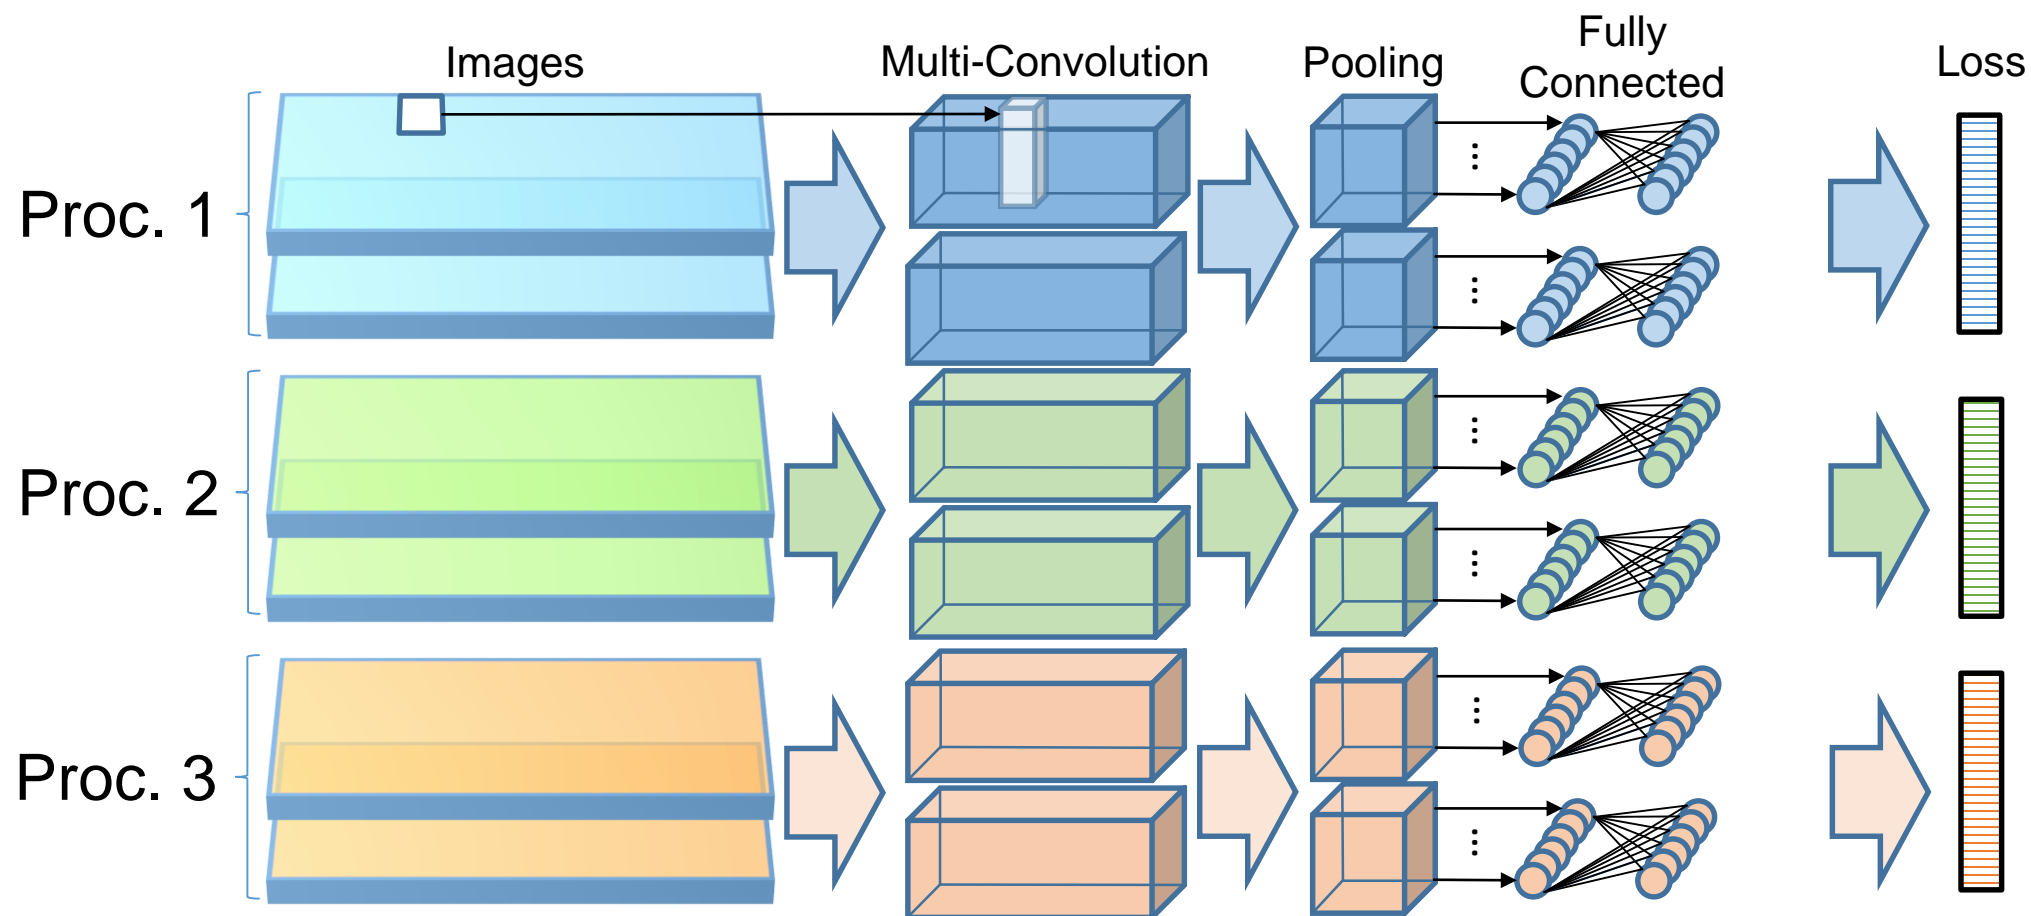
Parallelism



Data Parallelism




Data Parallelism



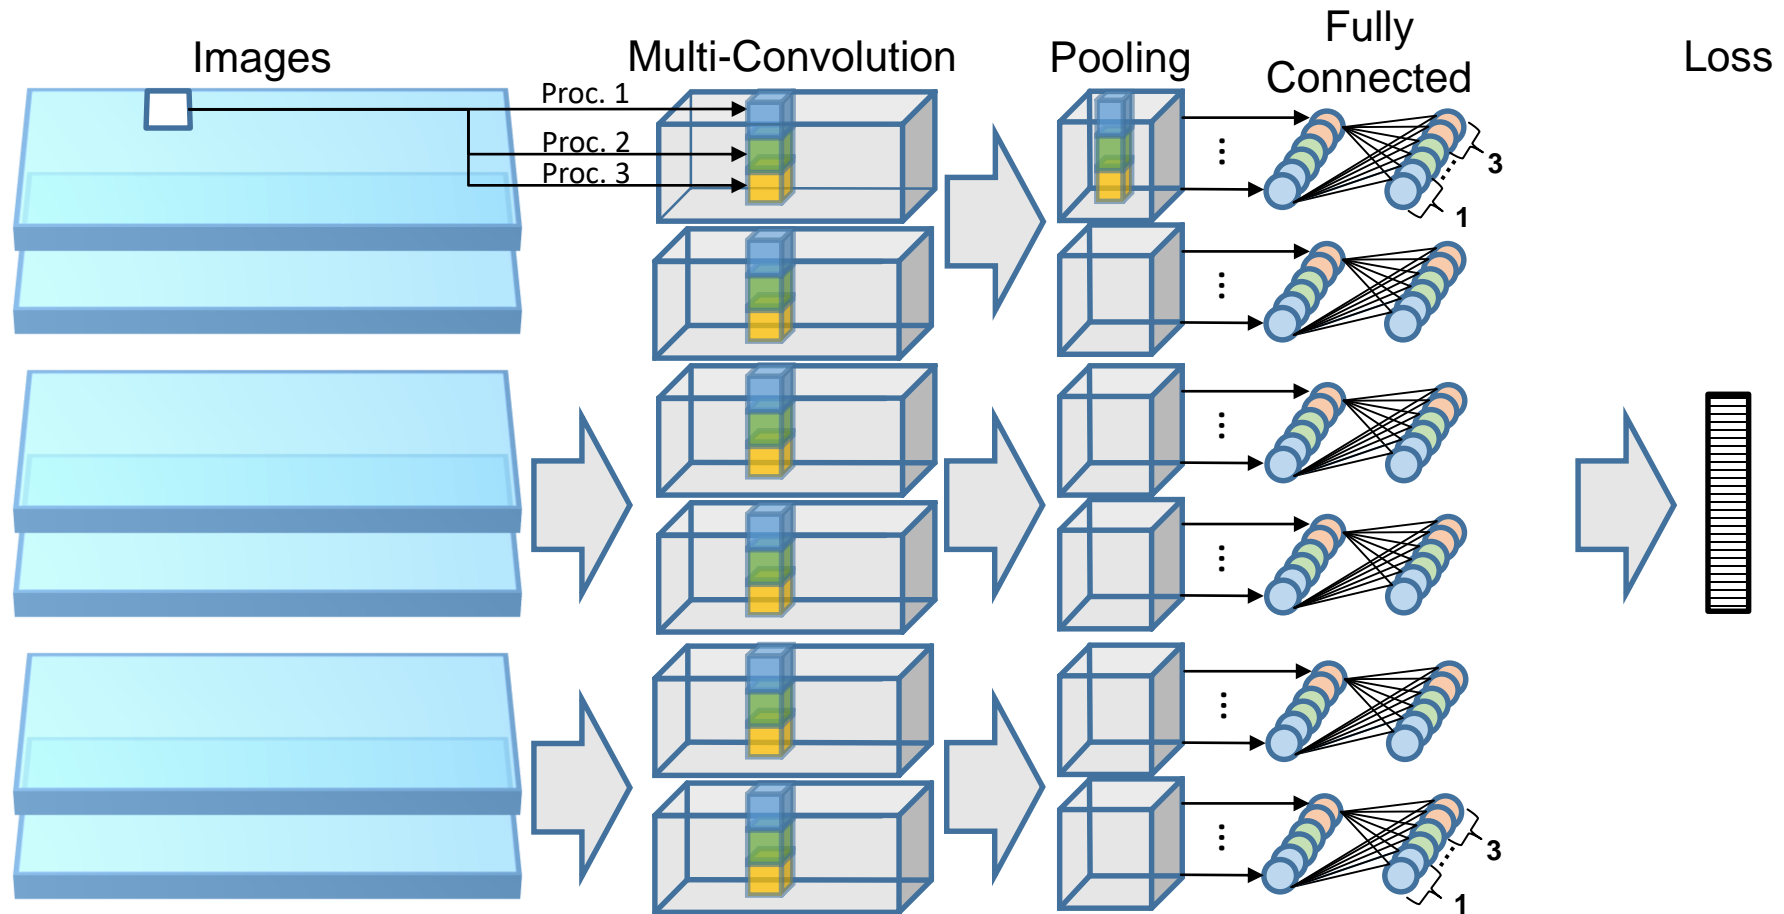
Data Parallelism

- ✓ Good for forward pass (independent)
- ✓ Backpropagation requires all-to-all communication only when accumulating results
- × Requires allocation of all parameters on each processor

Model Parallelism



Model Parallelism



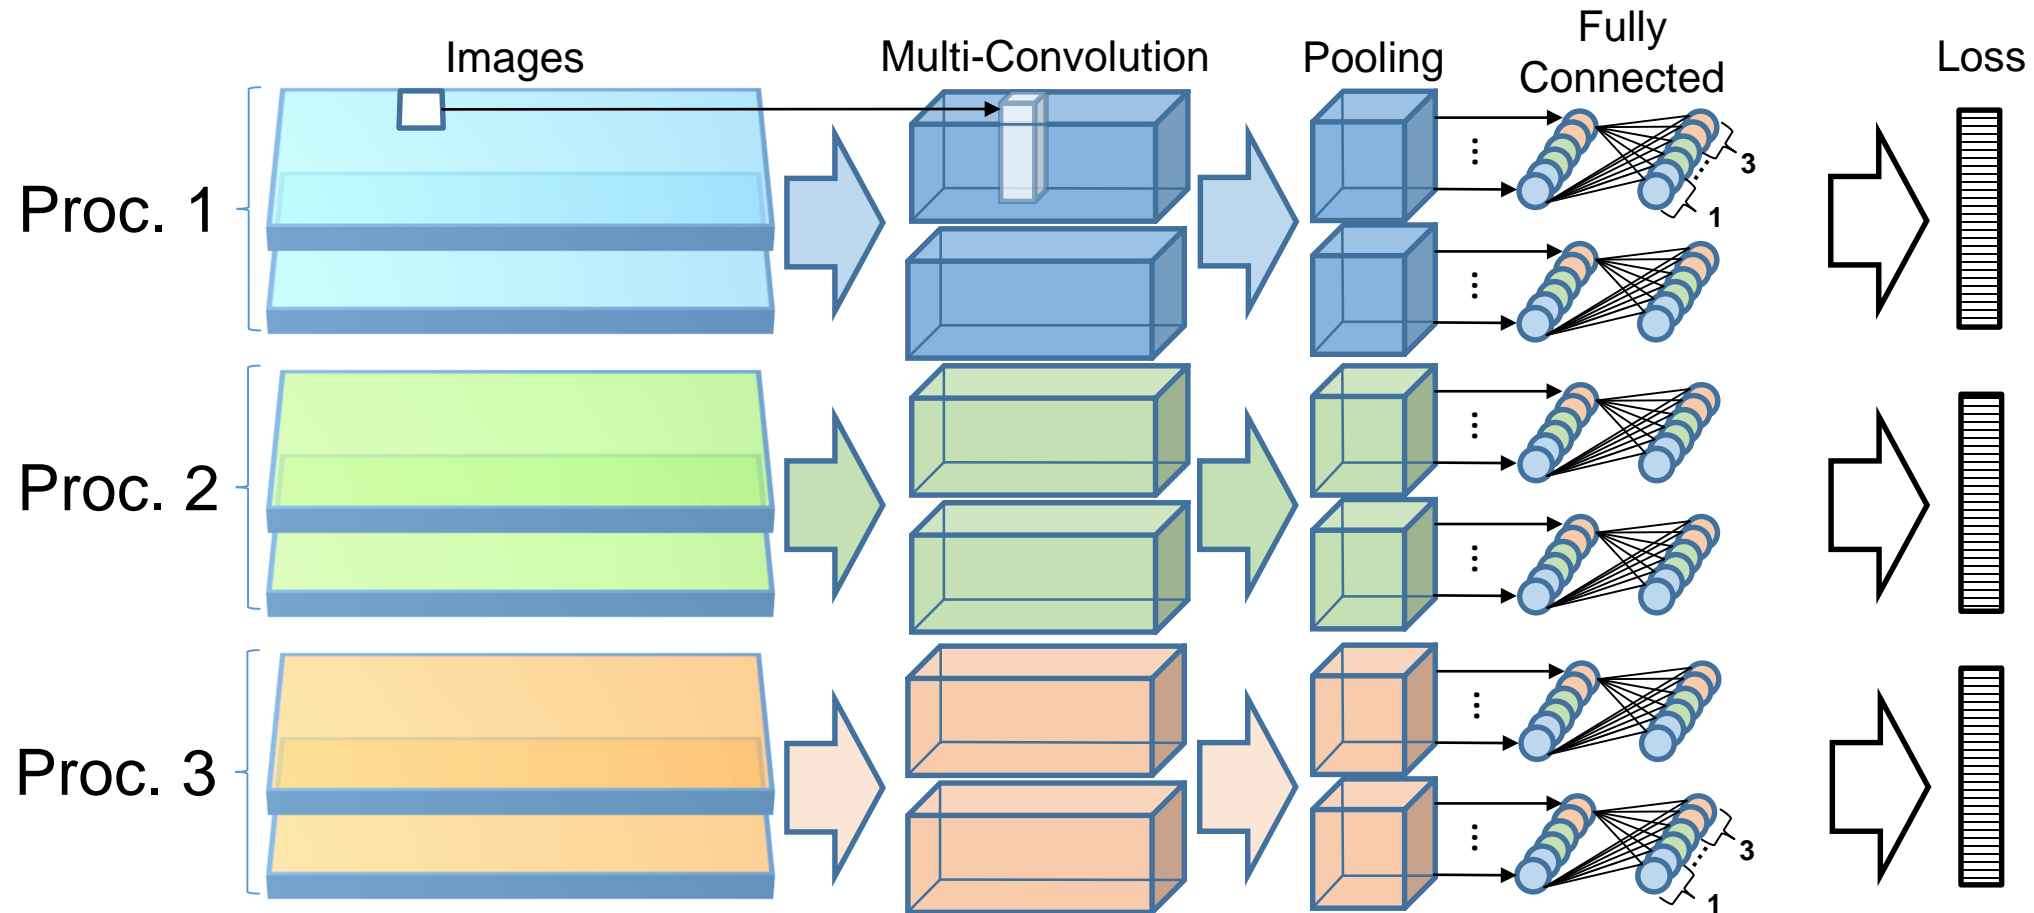
Model Parallelism

- ✓ Parameters can be divided across processors
- × Mini-batch has to be copied to all processors
- × Backpropagation requires all-to-all communication every layer

Hybrid Data/Model Parallelism

- Conjecture[Krizhevsky, 2014]: Most of the **computations** are performed in the convolutional portion, most of the **parameters** are stored in the fully connected portion
- Proposed Solution: Use data parallelism on convolutional portion and model parallelism on the FC portion

Hybrid Data/Model Parallelism [Krizhevsky, 2014]



[Krizhevsky. "One weird trick for parallelizing convolutional neural networks." \(2014\).](#)


Hybrid Data/Model Parallelism Results

- AlexNet, ILSVRC 2012:

GPUs	Batch size	Top-1 error	Time	Speedup
1	(128, 128)	42.33%	98.05h	1x
2	(256, 256)	42.63%	50.24h	1.95x
2	(256, 128)	42.27%	50.90h	1.93x
4	(512, 512)	42.59%	26.20h	3.74x
4	(512, 128)	42.44%	26.78h	3.66x
8	(1024, 1024)	43.28%	15.68h	6.25x
8	(1024, 128)	42.86%	15.91h	6.16x


[Krizhevsky. "One weird trick for parallelizing convolutional neural networks." \(2014\).](#)

Distributed Computing



Distributed Deep Learning

- Runs on a computer cluster
- Each node runs partially autonomously
- Inter-node communication from time to time
- Best result is gathered from the nodes
- Training data can be split to per-node “shards”




Distributed Deep Learning – Opportunities

- Increased memory:
 - More data
 - More parameters


- Fault tolerance
 - Protection against node crashes

- Improved stochasticity




Distributed Deep Learning – Determining Factors

- Computational independence
- Communication efficiency
- Network congestion
- Load balancing
- Points of failure



Distributed Synchronous SGD


- Communication step is added to the algorithm:
 - Initialize neural network weights (W_0)
 - For t in iterations:
 - Sample b images from dataset (B)
 - Compute loss $L_t(W_{t-1}, B)$
 - **Synchronize weights across workers**
 - Update weights using gradients and update rule g
- The step requires all nodes to have the same data ($\sum \nabla W$)
 - This collective operation is also called AllReduce
- Different ways to implement, depending on message size and network topology



Source: [Basics of Message-Passing](#)

Distributed Deep Learning – DistBelief

- Distributed learning infrastructure used at Google [Dean et al., 2012]
- Each model **replica** has the same parameters, but optimizes different data
 - Replicas are divided among several machines
- Two distributed optimization schemes for training:
 - Online – Downpour SGD
 - Batch – Sandblaster LBFGS
- Uses a centralized parameter server (several machines, sharded)
- Handles slow and faulty replicas



Asynchronous SGD – HOGWILD!

- To achieve coherency in distributed SGD, nodes must synchronize w.r.t. parameters:


- Each thread draws a random example i from the training data.
 - Acquire a lock on the current state of parameters θ .
 - Thread reads θ .
 - Thread updates $\theta \leftarrow (\theta - \alpha \nabla L(f_\theta(x_i), y_i))$.
 - Release lock on θ .

- *HOGWILD!* [Niu et al., 2011] removes this synchronization:

- Each thread draws a random example i from the training data.
- Thread reads current state of θ .
 - Thread updates $\theta \leftarrow (\theta - \alpha \nabla L(f_\theta(x_i), y_i))$.

Asynchronous SGD – HOGWILD!

- *HOGWILD!*:
 - Proven to converge in sparse problems
 - Provides near-linear scaling
 - Assumes shared-memory architecture (e.g., multicore CPUs)
- Formulates ML problems as hypergraphs $G = (V, E)$ where:
 - $w^* = \operatorname{argmin}_w f(w) = \operatorname{argmin}_w \sum_{e \in E} f_e(w_e)$
 - Each hyperedge $e \in E$ represents subsets of $[n]$
 - w_e is reduced to coordinates in e



Source: Wikipedia


Algorithm 1 HOGWILD! update for individual processors

```

1: loop
2:   Sample  $e$  uniformly at random from  $E$ 
3:   Read current state  $x_e$  and evaluate  $G_e(x_e)$ 
4:   for  $v \in e$  do  $x_v \leftarrow x_v - \gamma G_{ev}(x_e)$ 
5: end loop
  
```


Distributed Deep Learning – Downpour SGD

- Relaxation of *HOGWILD!* for distributed systems
- Algorithm:
 - Divide training data into subsets and run a replica on each subset
 - Every n_{fetch} iterations, fetch up-to-date parameters from server
 - Every n_{push} iterations, push local gradients to server
- Note that parameter shards may be “out-of-sync”




Distributed Deep Learning – Sandblaster LBFGS



- Coordinator process issues commands (dot product, scaling, multiplication, etc.) to slave nodes, each processing a different parameter shard
- Communication is sparser
 - Most of the information is stored locally
 - Coordinator messages are small
 - Slaves fetch parameters at the beginning of each batch, send gradients once in a while for fault tolerance
- Employs computation replication and load balancing
 - Nodes that finish their job get more jobs
 - If one node is slow, additional nodes get the same job



DistBelief Results – Time



DistBelief Results – Accuracy

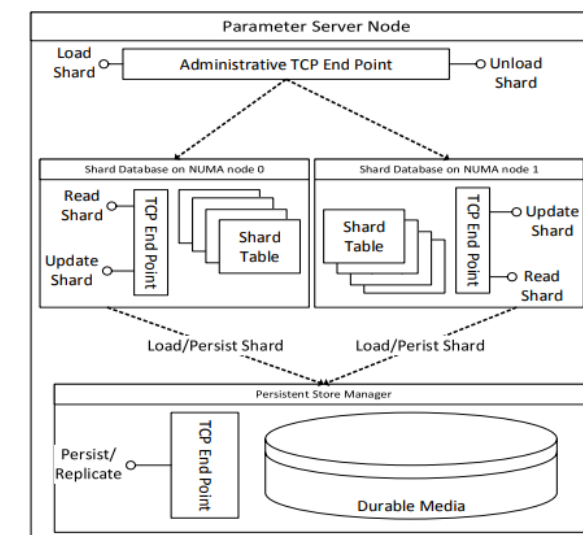
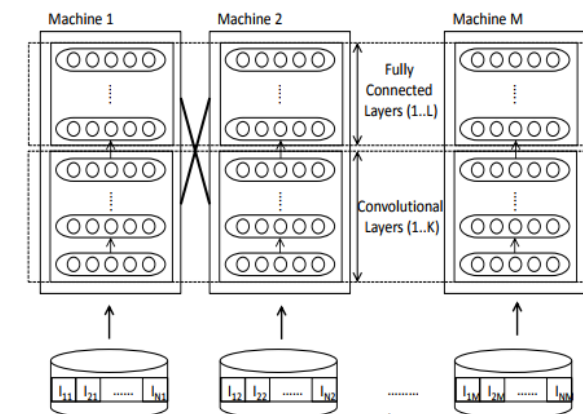


Project Adam


- Extends DistBelief with system-level support
 - Fast data serving mechanisms (e.g., with augmentation)
 - Better heterogeneous system management
 - Parameter server node optimization

- Bottom-up communication redesign
 - Control message, data message separation
 - Inter-node communication reduction
 - Weight differences are sent instead of weights

- Only system to train ImageNet22k



Hardware



Specialized Hardware


- GPU
 - Thousands of cores, massively parallel (5-14 TFLOP/s per card)
 - Multi-GPU nodes further increase training performance (using data/model parallelism)
 - Drawback: Hard to program efficiently. Solution: Specialized libraries (CUDNN)

- FPGA
 - Specialized for certain operations (e.g. convolutions)
 - Drawbacks: Even harder to program

- Convolutional Processing Units

Deep Learning with GPUs


- A distributed GPU-based system[Coates et al., 2013] was shown to run DistBelief-scale problems (1000 machines) with 3 multi-GPU nodes
- 3 tiers of concurrency: GPU, model parallelism, nodes



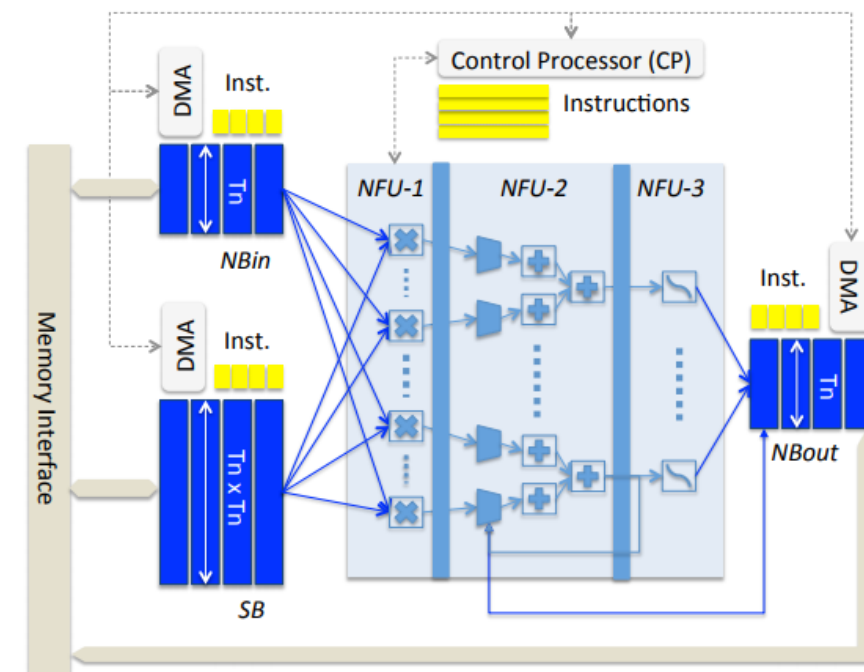
Time for a single mini-batch gradient update of a sparse autoencoder

Specialized Hardware

- Two approaches:



Mapping neurons to hardware



Custom processing elements

Specialized Hardware – ANNA



■ [Boser, Säckinger, Bromley, LeCun, Jackel, IEEE J. SSC 26(12), 1991]

■ 4096 Multiply-Accumulate operators

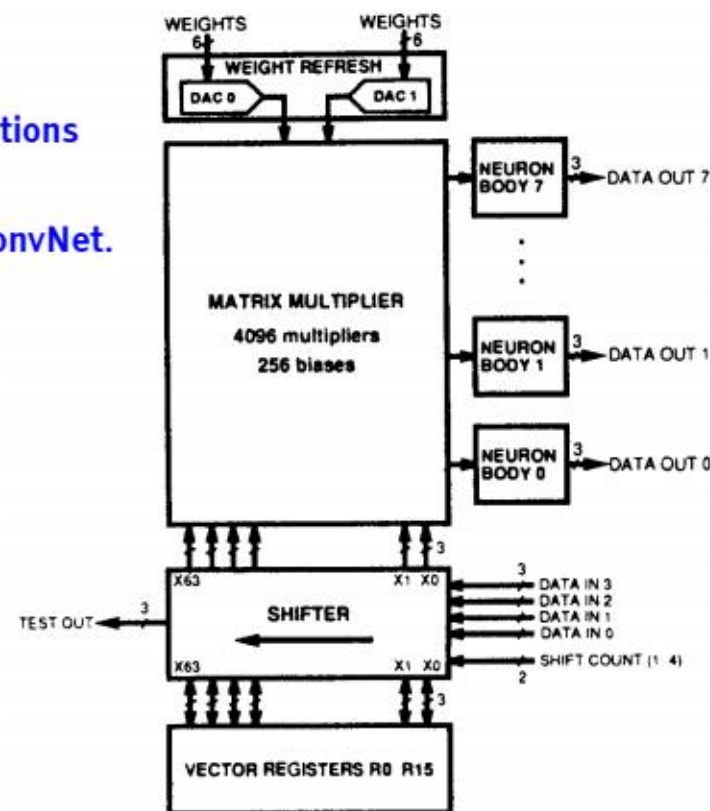

■ 6 bit weights, 4 bit states

■ 20 MHz clock

■ Shift registers for efficient I/O with convolutions

■ 4 GOPS (peak)

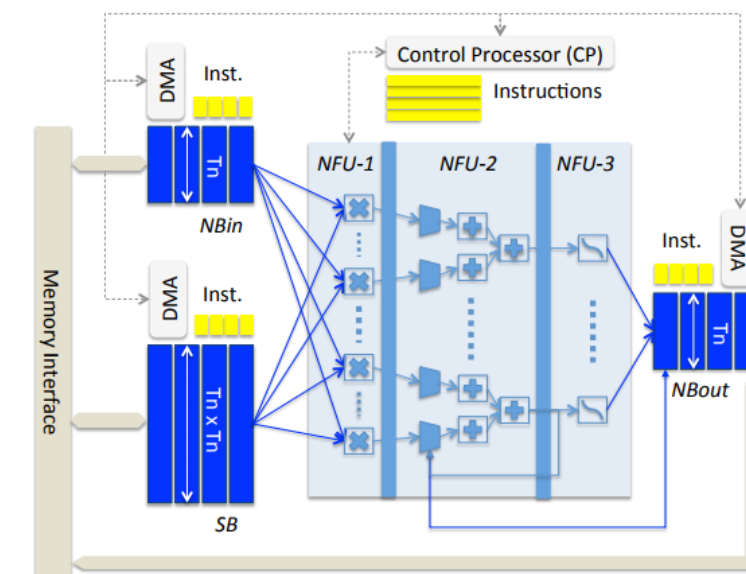

■ 1000 characters per second for OCR with ConvNet.



Source: Yann LeCun

Specialized Hardware – DianNao

- Instead of building the neural network on a circuit, creates a Neural Function Unit (NFU)
- Operations reflect the different layers
 - Computation
 - Reduction
 - Activation
- Each operation (conv., pooling, FC) takes up to three stages at computations (or less)
 - Computation
 - Reduction
 - Activation



Conclusions

- Many acceleration opportunities
- Architectures keep changing, and with them new techniques arise
- Algorithms can be modified (to some extent)
 - Proven “shortcuts” can be taken



Questions?